

Week 08: Graph Algorithms 2

Weighted Graphs

Weighted Graphs

2/61

Graphs so far have considered

- edge = an association between two vertices/nodes
- may be a precedence in the association (directed)

Some applications require us to consider


- a *cost* or *weight* of an association
- modelled by assigning values to edges (e.g. positive reals)

Weights can be used in both directed and undirected graphs.

... Weighted Graphs

3/61

Example: major airline flight routes in Australia



Representation: edge = direct flight; weight = approx flying time (hours)

... Weighted Graphs

4/61

Weights lead to minimisation-type questions, e.g.

1. Cheapest way to connect all vertices?

- a.k.a. *minimum spanning tree* problem
- assumes: edges are weighted and undirected

2. Cheapest way to get from A to B?

- a.k.a. *shortest path* problem
 - assumes: edge weights positive, directed or undirected
-

Exercise #1: Implementing a Route Finder

5/61

If we represent a street map as a graph

- what are the vertices?

- what are the edges?
- are edges directional?
- what are the weights?
- are the weights fixed?

What kind of algorithm would ...

- help us find the "quickest" way to get from A to B?

Weighted Graph Representation

6/61

Weights can easily be added to:

- adjacency matrix representation (0/1 \rightarrow int or float)
- adjacency lists representation (add int/float to list node)

An alternative representation useful in this context:


- edge list representation (list of (s,t,w) triples)

All representations work whether edges are directed or not.

... Weighted Graph Representation

7/61

Adjacency matrix representation with weights:




Note: need distinguished value to indicate "no edge".

... Weighted Graph Representation

8/61

Adjacency lists representation with weights:




Note: if undirected, each edge appears twice with same weight

... Weighted Graph Representation

9/61

Edge array / edge list representation with weights:



Note: not very efficient for use in processing algorithms, but does give a possible representation for min spanning trees or shortest paths

... Weighted Graph Representation

10/61

Sample adjacency matrix implementation in C requires minimal changes to previous Graph ADT:

WGraph.h

```
// edges are pairs of vertices (end-points) plus positive weight
typedef struct Edge {
    Vertex v;
    Vertex w;
    int    weight;
} Edge;

// returns weight, or 0 if vertices not adjacent
int adjacent(Graph, Vertex, Vertex);
```

... Weighted Graph Representation

11/61

WGraph.c

```
typedef struct GraphRep {
    int **edges; // adjacency matrix storing positive weights
                // 0 if nodes not adjacent
    int  nV;     // #vertices
    int  nE;     // #edges
} GraphRep;

void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));
    if (g->edges[e.v][e.w] == 0) { // edge e not in graph
        g->edges[e.v][e.w] = e.weight;
        g->nE++;
    }
}


int adjacent(Graph g, Vertex v, Vertex w) {
    assert(g != NULL && validV(g,v) && validV(g,w));
    return g->edges[v][w];
}
```

Minimum Spanning Trees

Exercise #2: Minimising Wires in Circuits

13/61

Electronic circuit designs often need to make the pins of several components electrically equivalent by wiring them together.



To interconnect a set of n pins we can use an arrangement of $n-1$ wires each connecting two pins.

What kind of algorithm would ...

- help us find the arrangement with the least amount of wire?

Minimum Spanning Trees

14/61

Reminder: *Spanning tree* ST of graph $G=(V,E)$

- *spanning* = all vertices, *tree* = no cycles
- ST is a subgraph of G ($G'=(V,E')$ where $E' \subseteq E$)
- ST is *connected* and *acyclic*

Minimum spanning tree MST of graph G

- MST is a spanning tree of G
- sum of edge weights is no larger than any other ST

Applications: Computer networks, Electrical grids, Transportation networks ...


Problem: how to (efficiently) find MST for graph G ?

NB: MST may not be unique (e.g. all edges have same weight \Rightarrow every ST is MST)


... Minimum Spanning Trees

15/61

Example:



An MST ...



... Minimum Spanning Trees

16/61

Brute force solution:

```
findMST(G):
  Input   graph G
  Output  a minimum spanning tree of G

  bestCost = ∞
  for all spanning trees t of G do
    if cost(t) < bestCost then
      bestTree = t
      bestCost = cost(t)
    end if
  end for
  return bestTree
```

Example of *generate-and-test* algorithm.

Not useful because #spanning trees is potentially large (e.g. n^{n-2} for a complete graph with n vertices)

... Minimum Spanning Trees

17/61

Simplifying assumption:

- edges in G are not directed (MST for digraphs is harder)

Kruskal's Algorithm

18/61

One approach to computing MST for graph G with V nodes:

1. start with empty MST
2. consider edges in increasing weight order
 - add edge if it does not form a cycle in MST
3. repeat until $V-1$ edges are added


Critical operations:

- iterating over edges in weight order
- checking for cycles in a graph

... Kruskal's Algorithm

19/61


Execution trace of Kruskal's algorithm:




Exercise #3: Kruskal's Algorithm

20/61


Show how Kruskal's algorithm produces an MST on:




After 3rd iteration:




After 6th iteration:



After 7th iteration:



After 8th iteration ($V-1=8$ edges added):



Pseudocode:

```
KruskalMST(G):  
  Input graph G with n nodes  
  Output a minimum spanning tree of G  
  
  MST=empty graph  
  sort edges(G) by weight  
  for each e in sortedEdgeList do  
    MST = MST  $\cup$  {e}  
    if MST has a cycle then  
      MST = MST  $\setminus$  {e}  
    end if  
    if MST has n-1 edges then  
      return MST  
    end if  
  end for
```

... Kruskal's Algorithm

23/61

Rough time complexity analysis ...

- sorting edge list is $O(E \cdot \log E)$
- at least V iterations over sorted edges
- on each iteration ...
 - getting next lowest cost edge is $O(1)$
 - checking whether adding it forms a cycle: cost = ??

Possibilities for cycle checking:

- use DFS ... too expensive?
- could use *Union-Find data structure* (see Sedgewick Ch.1)

Prim's Algorithm

24/61

Another approach to computing MST for graph $G=(V,E)$:

1. start from any vertex v and empty MST
2. choose edge not already in MST to add to MST
 - must be incident on a vertex s already connected to v in MST
 - must be incident on a vertex t not already connected to v in MST
 - must have minimal weight of all such edges
3. repeat until MST covers all vertices


Critical operations:

- checking for vertex being connected in a graph
- finding min weight edge in a set of edges

... Prim's Algorithm

25/61


Execution trace of Prim's algorithm (starting at $s=0$):



Exercise #4: Prim's Algorithm


26/61

Show how Prim's algorithm produces an MST on:




Start from vertex 0


After 1st iteration:




After 2nd iteration:




After 3rd iteration:



After 4th iteration:



After 8th iteration (all vertices covered):



Pseudocode:

```
PrimMST(G):  
  Input  graph G with n nodes  
  Output a minimum spanning tree of G  
  
  MST=empty graph  
  usedV={0}  
  unusedE=edges(g)  
  while |usedV|<n do  
    find e=(s,t,w)∈unusedE such that {  
      s∈usedV, t∉usedV and w is min weight of all such edges  
    }  
    MST = MST ∪ {e}  
    usedV = usedV ∪ {t}  
    unusedE = unusedE \ {e}  
  end while  
  return MST
```

Critical operation: finding best edge

Rough time complexity analysis ...

- V iterations of outer loop
 - in each iteration ...
 - find min edge with set of edges is $O(E) \Rightarrow O(V \cdot E)$ overall
 - find min edge with *priority queue* is $O(\log E) \Rightarrow O(V \cdot \log E)$ overall
-

Sidetrack: Priority Queues

Some applications of queues require

- items processed in order of "priority"
- rather than in order of entry (FIFO — first in, first out)

Priority Queues (PQueues) provide this via:

- **join**: insert item into PQueue with an associated priority (replacing enqueue)
- **leave**: remove item with highest priority (replacing dequeue)

Time complexity for naive implementation of a PQueue containing N items ...

- $O(1)$ for join $O(N)$ for leave

Most efficient implementation ("heap") ...

- $O(\log N)$ for join, leave
-

Other MST Algorithms

Boruvka's algorithm ... complexity $O(E \cdot \log V)$

- the oldest MST algorithm
- start with V separate components
- join components using min cost links
- continue until only a single component

Karger, Klein, and Tarjan ... complexity $O(E)$

- based on Boruvka, but non-deterministic
- randomly selects subset of edges to consider
- for the keen, here's [the paper](#) describing the algorithm

Shortest Path

Shortest Path

33/61

Path = sequence of edges in graph G $p = (v_0, v_1), (v_1, v_2), \dots, (v_{m-1}, v_m)$

cost(path) = sum of edge weights along path

Shortest path between vertices s and t

- a simple path $p(s, t)$ where $s = \text{first}(p)$, $t = \text{last}(p)$
- no other simple path $q(s, t)$ has $\text{cost}(q) < \text{cost}(p)$

Assumptions: weighted digraph, no negative weights.

Finding shortest path between two given nodes known as *source-target* SP problem

Variations: *single-source* SP, *all-pairs* SP

Applications: navigation, routing in data networks, ...

Single-source Shortest Path (SSSP)

34/61

Given: weighted digraph G , source vertex s

Result: shortest paths from s to all other vertices

- `dist[]` V -indexed array of cost of shortest path from s
- `pred[]` V -indexed array of predecessor in shortest path from s

Example:



Edge Relaxation


35/61

Assume: `dist[]` and `pred[]` as above (but containing data for shortest paths *discovered so far*)

`dist[v]` is length of shortest known path from s to v

`dist[w]` is length of shortest known path from s to w

Relaxation updates data for w if we find a shorter path from s to w :



Relaxation along edge $e=(v,w,\text{weight})$:

- if $\text{dist}[v] + \text{weight} < \text{dist}[w]$ then
 update $\text{dist}[w] := \text{dist}[v] + \text{weight}$ and $\text{pred}[w] := v$

Dijkstra's Algorithm

36/61

One approach to solving single-source shortest path problem ...

Data: $G, s, \text{dist}[], \text{pred}[]$ and

- $vSet$: set of vertices whose shortest path from s is unknown

Algorithm:

$\text{dist}[]$ // array of cost of shortest path from s
 $\text{pred}[]$ // array of predecessor in shortest path from s


```
dijkstraSSSP(G, source):
    Input graph G, source node

    initialise dist[] to all  $\infty$ , except  $\text{dist}[\text{source}] = 0$ 
    initialise pred[] to all -1
    vSet = all vertices of G
    while vSet  $\neq \emptyset$  do
        find  $s \in vSet$  with minimum  $\text{dist}[s]$ 
        for each  $(s, t, w) \in \text{edges}(G)$  do
            relax along  $(s, t, w)$ 
        end for
        vSet = vSet  $\setminus \{s\}$ 
    end while
```

Exercise #5: Dijkstra's Algorithm

37/61

Show how Dijkstra's algorithm runs on (source node = 0):



	[0]	[1]	[2]	[3]	[4]	[5]
dist	0	∞	∞	∞	∞	∞
pred	-	-	-	-	-	-

dist	0	14	9	7	∞	∞
pred	-	0	0	0	-	-

dist	0	14	9	7	∞	22
pred	-	0	0	0	-	3

dist	0	13	9	7	∞	12
pred	-	2	0	0	-	2

dist	0	13	9	7	20	12
pred	-	2	0	0	5	2

dist	0	13	9	7	18	12
pred	-	2	0	0	1	2

... Dijkstra's Algorithm

39/61

Why Dijkstra's algorithm is correct:

Hypothesis.

- (a) For visited $s \dots dist[s]$ is shortest distance from source
- (b) For unvisited $t \dots dist[t]$ is shortest distance from source *via visited nodes*

Proof.

Base case: no visited nodes, $dist[source]=0$, $dist[s]=\infty$ for all other nodes

Induction step:

1. If s is unvisited node with minimum $dist[s]$, then $dist[s]$ is shortest distance from source to s :
 - if \exists shorter path via only visited nodes, then $dist[s]$ would have been updated when processing the predecessor of s on this path
 - if \exists shorter path via an unvisited node u , then $dist[u] < dist[s]$, which is impossible if s has min distance of all unvisited nodes
2. This implies that (a) holds for s after processing s
3. (b) still holds for all unvisited nodes t after processing s :
 - if \exists shorter path via s we would have just updated $dist[t]$
 - if \exists shorter path without s we would have found it previously

... Dijkstra's Algorithm

40/61

Time complexity analysis ...

Each edge needs to be considered once $\Rightarrow O(E)$.

Outer loop has $O(V)$ iterations.

Implementing "**find** $s \in vSet$ **with** minimum $dist[s]$ "

1. try all $s \in vSet \Rightarrow cost = O(V) \Rightarrow overall cost = O(E + V^2) = O(V^2)$
2. using a PQueue to implement extracting minimum
 - can improve overall cost to $O(E + V \cdot \log V)$ (for best-known implementation)

41/61


All-pair Shortest Path (APSP)

Given: weighted digraph G

Result: shortest paths between all pairs of vertices

- $\text{dist}[\][\]$ $V \times V$ -indexed matrix of cost of shortest path from v_{row} to v_{col}
- $\text{path}[\][\]$ $V \times V$ -indexed matrix of next node in shortest path from v_{row} to v_{col}

Example:



Weighted Digraph

V	0	1	2	3	4	
0	0	0.3	0.7	1.1	inf	dist
1	1.8	0	0.6	1.0	inf	
2	1.2	0.5	0	0.4	inf	
3	0.8	1.1	1.5	0	inf	
4	1.3	0.6	0.1	0.5	0	
0	-	1	2	2	-	path
1	2	-	2	2	-	
2	3	1	-	3	-	
3	0	0	0	-	-	
4	2	2	2	2	-	

Shortest paths between all vertices

Floyd's Algorithm

42/61

One approach to solving all-pair shortest path problem...

Data: $G, \text{dist}[\][\], \text{path}[\][\]$ Algorithm:

```
dist[][] // array of cost of shortest path from s to t
path[][] // array of next node after s on shortest path from s to t
```


```
floydAPSP(G):
    Input graph G

    initialise dist[s][t]=0 for each s=t
                          =w for each (s,t,w)∈edges(G)
                          =∞ otherwise
    initialise path[s][t]=t for each (s,t,w)∈edges(G), otherwise to -1
    for all i∈vertices(G) do
        for all s∈vertices(G) do
            for all t∈vertices(G) do
                if dist[s][i]+dist[i][t] < dist[s][t] then
                    dist[s][t]=dist[s][i]+dist[i][t]
                    path[s][t]=path[s][i]
                end if
            end for
        end for
    end for
```

Exercise #6: Floyd's Algorithm

43/61

Show how Floyd's algorithm runs on:



After 1st iteration i=0: unchanged

After 2nd iteration i=1:

dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	14	9	7	∞	∞	[0]	–	1	2	3	–	–
[1]	∞	0	∞	∞	5	∞	[1]	–	–	–	–	4	–
[2]	∞	4	0	∞	9	3	[2]	–	1	–	–	1	5
[3]	∞	∞	10	0	∞	15	[3]	–	–	2	–	–	5
[4]	∞	∞	∞	∞	0	∞	[4]	–	–	–	–	–	–
[5]	∞	∞	∞	∞	2	0	[5]	–	–	–	–	4	–

After 3rd iteration i=2:

dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	13	9	7	18	12	[0]	–	2	2	3	2	2
[1]	∞	0	∞	∞	5	∞	[1]	–	–	–	–	4	–
[2]	∞	4	0	∞	9	3	[2]	–	1	–	–	1	5
[3]	∞	14	10	0	19	13	[3]	–	2	2	–	2	2
[4]	∞	∞	∞	∞	0	∞	[4]	–	–	–	–	–	–
[5]	∞	∞	∞	∞	2	0	[5]	–	–	–	–	4	–

After 4th iteration i=3: unchanged

After 5th iteration i=4: unchanged


After 6th iteration i=5:

dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	13	9	7	14	12	[0]	–	2	2	3	2	2
[1]	∞	0	∞	∞	5	∞	[1]	–	–	–	–	4	–
[2]	∞	4	0	∞	5	3	[2]	–	1	–	–	5	5
[3]	∞	14	10	0	15	13	[3]	–	2	2	–	2	2
[4]	∞	∞	∞	∞	0	∞	[4]	–	–	–	–	–	–
[5]	∞	∞	∞	∞	2	0	[5]	–	–	–	–	4	–

Why Floyd's algorithm is correct:

A shortest path from s to t using only nodes from $\{0, \dots, i\}$ is the shorter of

- a shortest path from s to t using only nodes from $\{0, \dots, i-1\}$
- a shortest path from s to i using only nodes from $\{0, \dots, i-1\}$ plus a shortest path from i to t using only nodes from $\{0, \dots, i-1\}$



Also known as Floyd-Warshall algorithm (can you see why?)

... Floyd's Algorithm

Cost analysis ...

- initialising $dist[][], path[][] \Rightarrow O(E)$
- V iterations to update $dist[][], path[][] \Rightarrow O(V^3)$


Time complexity of Floyd's algorithm: $O(V^3)$ (same as Warshall's algorithm for transitive closure)

Network Flow

Exercise #7: Merchandise Distribution

Lucky Cricket Company produces cricket balls in Fairfield and has a warehouse in Rozelle that stocks them.

To ship the cricket balls from the factory to the warehouse, Lucky Cricket leases space on trucks that have limited capacity:



What kind of algorithm would ...

- help us find the maximum number of crates that can be shipped from Fairfield to Rozelle per day?

Flow Networks

Flow network ...

- weighted graph $G=(V,E)$
- distinct nodes $s \in V$ (source), $t \in V$ (sink)

Edge weights denote *capacities*

Applications:

- Distribution networks, e.g.
 - source: oil field
 - sink: refinery
 - edges: pipes
- Traffic flow

... Flow Networks

50/61

Flow in a network $G=(V,E)$... nonnegative $f(v,w)$ for all vertices $v,w \in V$ such that

- $f(v,w) \leq \text{capacity}$ for each edge $e=(v,w,\text{capacity}) \in E$
- $f(v,w)=0$ if no edge between v and w
- total flow into a vertex = total flow out of a vertex:


$$\sum_{x \in V} f(x, v) = \sum_{y \in V} f(v, y) \quad \text{for all } v \in V \setminus \{s, t\}$$

Maximum flow ... no other flow from s to t has larger value


... Flow Networks

51/61

Example:



A (maximum) flow ...




Augmenting Paths

52/61

Assume ... $f(v,w)$ contains current flow

Augmenting path: any path from source s to sink t that can currently take more flow

Example:



Residual Network


53/61

Assume ... flow network $G=(V,E)$ and flow $f(v,w)$

Residual network (V, E') :

- same vertex set V
- for each edge $v \rightarrow^c w \in E \dots$
 - $f(v, w) < c \Rightarrow$ add edge $(v \rightarrow^{c-f(v, w)} w)$ to E'
 - $f(v, w) > 0 \Rightarrow$ add edge $(v \leftarrow^{f(v, w)} w)$ to E'


Example:



Exercise #8: Augmenting Paths and Residual Networks


54/61

Find an augmenting path in:




and show the residual network after augmenting the flow

1. Augmenting path:



maximum additional flow = 1

2. Residual network:



Can you find a further augmenting path in the new residual network?

Edmonds-Karp Algorithm

56/61

One approach to solving maximum flow problem ...

maxflow(G):

1. Find a shortest augmenting path
2. Update flow[][] so as to represent residual graph
3. Repeat until no augmenting path can be found

... Edmonds-Karp Algorithm

57/61

Algorithm:

```
flow[ ][ ] // V×V array of current flow
visited[ ] /* array of predecessor nodes on shortest path
            from source to sink in residual network */
```

```
maxflow(G):
    Input  flow network G with source s and sink t
    Output maximum flow value

    initialise flow[v][w]=0 for all vertices v, w
    maxflow=0
    while ∃shortest augmenting path visited[ ] from s to t do
        df = maximum additional flow via visited[ ]
        // adjust flow so as to represent residual graph
        v=t
        while v≠s do
            flow[visited[v]][v] = flow[visited[v]][v] + df;
            flow[v][visited[v]] = flow[v][visited[v]] - df;
            v=visited[v]
        end while
        maxflow=maxflow+df
    end while
    return maxflow
```

Shortest augmenting path can be found by standard BFS

... Edmonds-Karp Algorithm

58/61

Time complexity analysis ...

- *Theorem.* The number of augmenting paths needed is at most $V \cdot E/2$.
⇒ Outer loop has $O(V \cdot E)$ iterations.
- Finding augmenting path ⇒ $O(E)$.


Overall cost of Edmonds-Karp algorithm: $O(V \cdot E^2)$

Note: Edmonds-Karp algorithm is an implementation of general *Ford-Fulkerson method*

Exercise #9: Edmonds-Karp Algorithm

59/61

Show how Edmonds-Karp algorithm runs on:



flow	[0]	[1]	[2]	[3]	[4]	[5]	c-f	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	0	0	0	0	0	[0]	-	2	3	-	-	-
[1]	0	0	0	0	0	0	[1]	-	-	-	3	1	-
[2]	0	0	0	0	0	0	[2]	-	-	-	1	1	-
[3]	0	0	0	0	0	0	[3]	-	-	-	-	-	2
[4]	0	0	0	0	0	0	[4]	-	-	-	-	-	3
[5]	0	0	0	0	0	0	[5]	-	-	-	-	-	-

augmenting path: 0-1-3-5, df: 2

flow	[0]	[1]	[2]	[3]	[4]	[5]	c-f	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	2	0	0	0	0	[0]	-	0	3	-	-	-
[1]	-2	0	0	2	0	0	[1]	2	-	-	1	1	-
[2]	0	0	0	0	0	0	[2]	-	-	-	1	1	-
[3]	0	-2	0	0	0	2	[3]	-	2	-	-	-	0
[4]	0	0	0	0	0	0	[4]	-	-	-	-	-	3
[5]	0	0	0	-2	0	0	[5]	-	-	-	2	-	-

augmenting path: 0-2-4-5, df: 1

flow	[0]	[1]	[2]	[3]	[4]	[5]	c-f	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	2	1	0	0	0	[0]	-	0	2	-	-	-
[1]	-2	0	0	2	0	0	[1]	2	-	-	1	1	-
[2]	-1	0	0	0	1	0	[2]	1	-	-	1	0	-
[3]	0	-2	0	0	0	2	[3]	-	2	-	-	-	0
[4]	0	0	-1	0	0	1	[4]	-	-	1	-	-	2
[5]	0	0	0	-2	-1	0	[5]	-	-	-	2	1	-

augmenting path: 0-2-3-1-4-5, df: 1

flow	[0]	[1]	[2]	[3]	[4]	[5]	c-f	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	2	2	0	0	0	[0]	-	0	1	-	-	-
[1]	-2	0	0	1	1	0	[1]	2	-	-	1	0	-
[2]	-2	0	0	1	1	0	[2]	2	-	-	0	0	-
[3]	0	-1	-1	0	0	2	[3]	-	1	1	-	-	0
[4]	0	-1	-1	0	0	2	[4]	-	1	1	-	-	1
[5]	0	0	0	-2	-2	0	[5]	-	-	-	2	2	-

- Weighted graph representations
- Minimum Spanning Tree (MST)
 - Kruskal, Prim
- Shortest path problems
 - Dijkstra (single source SPP)
 - Floyd (all-pair SSP)
- Flow networks
 - Edmonds-Karp (maximum flow)

- Suggested reading (Sedgewick):
 - MST ... Ch.20-20.4
 - SSP ... Ch.21-21.3
 - Flow ... Ch.22.1-22.2