Week 09: Search Tree Algorithms 1

Searching

An extremely common application in computing

- given a (large) collection of *items* and a *key* value
- find the item(s) in the collection containing that key
    - item = (key, val$_1$, val$_2$, …) (i.e. a structured data type)
    - key = value used to distinguish items (e.g. student ID)

Applications: Google, databases, ...

... Searching

Since searching is a very important/frequent operation,
many approaches have been developed to do it

Linear structures: arrays, linked lists, files

Arrays = random access. Lists, files = sequential access.

Cost of searching:

	Array	List	File
Unsorted	O(n) (linear scan)	O(n) (linear scan)	O(n) (linear scan)
Sorted	O(log n) (binary search)	O(n) (linear scan)	O(log n) (*seek*, *seek>*, …)

- *O(n)* … linear scan   (search technique of last resort)
- *O(log n)* … binary search, *search trees*   (trees also have other uses)

Also (cf. COMP9021): hash tables   (*O(1)*, but only under optimal conditions)

### ... Searching

Maintaining the order in sorted arrays and files is a costly operation.

Search trees are as efficient to search but more efficient to maintain.

Example: the following tree corresponds to the sorted array `[2,5,10,12,14,17,20,24,29,30,31,32]`:

Tree Data Structures

Trees

Trees are connected graphs

- consisting of nodes and edges (called *links*), with no cycles (no "up-links")
- each node contains a data value   (or key+data)
- each node has links to $\leq k$ other child nodes   (*k=2* below)

... Trees

Trees are used in many contexts, e.g.

- representing hierarchical data structures (e.g. expressions)
- efficient searching (e.g. sets, symbol tables, …)

... Trees

Trees can be used as a data structure, but also for *illustration*.

E.g. showing evaluation of a prefix arithmetic expression

Binary trees (*k=2* children per node) can be defined recursively, as follows:

A *binary tree* is either

- empty   (contains no nodes)
- consists of a *node*, with two *subtrees*
  - node contains a value
  - left and right subtrees are *binary trees*

Other special kinds of tree

- *m-ary tree*: each internal node has exactly *m* children
- *Ordered tree*: all left values < root, all right values > root
- *Balanced tree*: has ≅minimal height for a given number of nodes
- *Degenerate tree*: has ≅maximal height for a given number of nodes

Search Trees

Binary Search Trees

*Binary search trees* (or *BSTs*) have the characteristic properties

- each node is the root of 0, 1 or 2 subtrees
- all values in any left subtree are less than root
- all values in any right subtree are greater than root
- these properties applies over all nodes in the tree

(perfectly) *balanced trees* have the properties

- #nodes in left subtree = #nodes in right subtree
- this property applies over all nodes in the tree

Three binary search trees

Operations on BSTs:

- *insert*(Tree,Item) … add new item to tree via key
- *delete*(Tree,Key) … remove item with specified key from tree
- *search*(Tree,Key) … find item containing key in tree
- plus, "bookkeeping" … new(), free(), show(), …

Notes:

- nodes contain `Items`; we just show `Item.key`
- keys are unique   (not technically necessary)

Examples of binary search trees:


Shape of tree is determined by order of insertion.

Level of node = path length from root to node

Height (or: *depth*) of tree = max path length from root to leaf

Height-balanced tree: $\forall$ nodes: height(left subtree) = height(right subtree)

Time complexity of tree algorithms is typically *O(height)*

Exercise #1: Insertion into BSTs

For each of the sequences below

- start from an initially empty binary search tree
- show tree resulting from inserting values in order given

(a) 4 2 6 5 1 7 3

(b) 6 5 2 3 4 7 1

(c) 1 2 3 4 5 6 7

Assume new values are always inserted as new leaf nodes

(a) the balanced tree from 3 slides ago (height = 2)

(b) the non-balanced tree from 3 slides ago (height = 4)

(c) a fully degenerate tree of height 6

Representing BSTs

Binary trees are typically represented by node structures

- containing a value, and pointers to child nodes

Most tree algorithms move *down* the tree.
If upward movement needed, add a pointer to parent.

Typical data structures for trees …

```
// a Tree is represented by a pointer to its root node
typedef struct Node *Tree;

// a Node contains its data, plus left and right subtrees
typedef struct Node {
   int  data;
   Tree left, right;
} Node;

// some macros that we will use frequently
#define data(tree)  ((tree)->data)
#define left(tree)  ((tree)->left)
#define right(tree) ((tree)->right)
```

We ignore items $\Rightarrow$ `data` in `Node` is just a key

Abstract data vs concrete data …

Tree Algorithms

Most tree algorithms are best described recursively:

```
TreeSearch(tree,item):
|   Input  tree, item
|   Output true if item found in tree, false otherwise
|
|   if tree is empty then
|      return false
|   else if item < data(tree) then
|      return TreeSearch(left(tree),item)
|   else if item > data(tree) then
|      return TreeSearch(right(tree),item)
|   else            // found
|     return true
|   end if
```

Insert an item into appropriate subtree:

```
insertAtLeaf(tree,item):
|   Input  tree, item
|   Output tree with item inserted
|
|   if tree is empty then
|      return new node containing item
|   else if item < data(tree) then
|      return insertAtLeaf(left(tree),item)
|   else if item > data(tree) then
|      return insertAtLeaf(right(tree),item)
|   else
|      return tree    // avoid duplicates
|   end if
```

Tree Traversal

Iteration (traversal) on …

- `Lists` … visit each value, from first to last
- `Graphs` … visit each vertex, order determined by DFS/BFS/…

For binary `Tree`s, several well-defined visiting orders exist:

- *preorder* (NLR) … visit root, then left subtree, then right subtree
- *inorder* (LNR) … visit left subtree, then root, then right subtree
- *postorder* (LRN) … visit left subtree, then right subtree, then root
- *level-order* … visit root, then all its children, then all their children

... Tree Traversal

Consider "visiting" an expression tree like:


NLR: + * 1 3 - * 5 7 9   (prefix-order: useful for building tree)
LNR: 1 * 3 + 5 * 7 - 9   (infix-order: "natural" order)
LRN: 1 3 * 5 7 * 9 - +   (postfix-order: useful for evaluation)
Level: + * - 1 3 * 9 5 7   (level-order: useful for printing tree)

Exercise #2: Tree Traversal

Show NLR, LNR, LRN traversals for the following tree:

NLR (preorder): 20 10 5 2 14 12 17 30 24 29 32 31

LNR (inorder): 2 5 10 12 14 17 20 24 29 30 31 32

LRN (postorder): 2 5 12 17 14 10 29 24 31 32 30 20

Exercise #3: Non-recursive traversals

Write a non-recursive *preorder* traversal algorithm.

Assume that you have a stack ADT available.

```
showBSTreePreorder(t):
|   Input tree t
|
|   push t onto new stack S
|   while stack is not empty do
|   |   t=pop(S)
|   |   print data(t)
|   |   if right(t) is not empty then
|   |       push right(t) onto S
|   |   end if
|   |   if left(t) is not empty then
|   |       push left(t) onto S
|   |   end if
|   end while
```

Joining Two Trees

An auxiliary tree operation …

Tree operations so far have involved just one tree.

An operation on two trees: `t = joinTrees(t`$_1$`,t`$_2$`)`

- Pre-conditions:
  - takes two BSTs; returns a single BST
  - $max(key(t_1)) < min(key(t_2))$
- Post-conditions:
  - result is a BST (i.e. fully ordered)
  - containing all items from t_1 and t_2

... Joining Two Trees

Method for performing tree-join:

- find the min node in the right subtree (t_2)
- replace min node by its right subtree
- elevate min node to be new root of both trees

Advantage: doesn't increase height of tree significantly

$x \leq \text{height}(t) \leq x+1$, where $x = \max(\text{height}(t_1), \text{height}(t_2))$

Variation: choose deeper subtree; take root from there.

Joining two trees:

Note: t2' may be less deep than t2

Implementation of tree-join:

```
joinTrees(t₁,t₂):
|   Input  trees t₁,t₂
|   Output t₁ and t₂ joined together
|
|   if t₁ is empty then return t₂
|   else if t₂ is empty then return t₁
|   else
|   |   curr=t₂, parent=NULL
|   |   while left(curr) is not empty do     // find min element in t₂
|   |       parent=curr
|   |       curr=left(curr)
|   |   end while
|   |   if parent≠NULL then
|   |       left(parent)=right(curr)  // unlink min element from parent
|   |       right(curr)=t₂
|   |   end if
|   |   left(curr)=t₁
|   |   return curr                    // curr is new root
|   end if
```

Join the trees

Deletion from BSTs

Insertion into a binary search tree is easy.

Deletion from a binary search tree is harder.

Four cases to consider …

- empty tree … new tree is also empty
- zero subtrees … unlink node from parent
- one subtree … replace by child
- two subtrees … replace by successor, join two subtrees

... Deletion from BSTs

Case 2: item to be deleted is a leaf (zero subtrees)

Just delete the item

... Deletion from BSTs

Case 2: item to be deleted is a leaf (zero subtrees)

deleted k ...

... Deletion from BSTs

Case 3: item to be deleted has one subtree

delete p ...

Replace the item by its only subtree

... Deletion from BSTs

Case 3: item to be deleted has one subtree

deleted p ...

... Deletion from BSTs

Case 4: item to be deleted has two subtrees

delete m ...

Version 1: right child becomes new root, attach left subtree to min element of right subtree

... Deletion from BSTs

Case 4: item to be deleted has two subtrees

deleted m (v1) ...

... Deletion from BSTs

Case 4: item to be deleted has two subtrees

delete m ...

Version 2: *join* left and right subtree

... Deletion from BSTs

Case 4: item to be deleted has two subtrees

deleted m (v2) ...

... Deletion from BSTs

Pseudocode (version 2):

```
TreeDelete(t,item):
|   Input  tree t, item
|   Output t with item deleted
|
|   if t is not empty then          // nothing to do if tree is empty
|   |   if item < data(t) then       // delete item in left subtree
|   |       left(t)=TreeDelete(left(t),item)
|   |   else if item > data(t) then  // delete item in left subtree
|   |       right(t)=TreeDelete(right(t),item)
```

```
│   │   else                      // node 't' must be deleted
│   │   │   if left(t) and right(t) are empty then
│   │   │       new=empty tree              // 0 children
│   │   │   else if left(t) is empty then
│   │   │       new=right(t)                // 1 child
│   │   │   else if right(t) is empty then
│   │   │       new=left(t)                 // 1 child
│   │   │   else
│   │   │       new=joinTrees(left(t),right(t))  // 2 children
│   │   │   end if
│   │   │   free memory allocated for t
│   │   │   t=new
│   │   end if
│   end if
│   return t
```

Balanced BSTs

Balanced Binary Search Trees

Goal: build binary search trees which have

- minimum height $\Rightarrow$ minimum worst case search cost

Perfectly balanced tree with *N* nodes has

- abs(#nodes(LeftSubtree) - #nodes(RightSubtree)) < 2, for every node
- height of $log_2 N \Rightarrow$ worst case search *O(log N)*

Three *strategies* to improving worst case search in BSTs:

- *randomise* — reduce chance of worst-case scenario occuring
- *amortise* — do more work at insertion to make search faster
- *optimise* — implement all operations with performance bounds

Operations for Rebalancing

To assist with rebalancing, we consider new operations:

Left rotation

- move right child to root; rearrange links to retain order

Right rotation

- move left child to root; rearrange links to retain order

Insertion at root

- each new item is added as the new root node

Tree Rotation

In tree below: $t_1 < n_2 < t_2 < n_1 < t_3$

... Tree Rotation

Method for rotating tree T right:

- N_1 is current root; N_2 is root of N_1's left subtree
- N_1 gets new left subtree, which is N_2's right subtree
- N_1 becomes root of N_2's new right subtree
- N_2 becomes new root

Left rotation: swap left/right in the above.

Cost of tree rotation: *O(1)*

... Tree Rotation

Algorithm for right rotation:

```
rotateRight(n₁):
|   Input  tree n₁
|   Output n₁ rotated to the right
|
|   if n₁ is empty or left(n₁) is empty then
|       return n₁
|   end if
|   n₂=left(n₁)
|   left(n₁)=right(n₂)
|   right(n₂)=n₁
|   return n₂
```


## Exercise #5: Tree Rotation

Consider the tree `t`:

Show the result of `rotateRight(t)`

Exercise #6: Tree Rotation

Write the algorithm for left rotation


```
rotateLeft(n₂):
|  Input  tree n₂
|  Output n₂ rotated to the left
|
|  if n₂ is empty or right(n₂) is empty then
|      return n₂
|  end if
|  n₁=right(n₂)
|  right(n₂)=left(n₁)
|  left(n₁)=n₂
|  return n₁
```

Insertion at Root

Previous description of BSTs inserted at leaves.

Different approach: insert new item at root.

Potential disadvantages:

- large-scale rearrangement of tree for each insert

Potential advantages:

- recently-inserted items are close to root
- low cost if recent items more likely to be searched

... Insertion at Root

Method for inserting at root:

- base case:
  - tree is empty; make new node and make it root
- recursive case:
  - insert new node as root of appropriate subtree
 - lift new node to root by rotation

... Insertion at Root

Exercise #7: Insertion at Root

Consider the tree `t`:

Show the result of `insertAtRoot(t,24)`

... Insertion at Root

Analysis of insertion-at-root:

- same complexity as for insertion-at-leaf: *O(height)*
- tendency to be balanced, but no balance guarantee
- benefit comes in searching
    - for some applications, search favours recently-added items
    - insertion-at-root ensures these are close to root
- could even consider "move to root when found"
    - effectively provides "self-tuning" search tree

Rebalancing Trees

An approach to balanced trees:

- insert into leaves as for simple BST
- periodically, rebalance the tree

Question: how frequently/when/how to rebalance?

```
NewTreeInsert(tree,item):
|   Input  tree, item
|   Output tree with item randomly inserted
|
|   t=insertAtLeaf(tree,item)
|   if #nodes(t) mod k = 0 then
|      t=rebalance(t)
|   end if
|   return t
```

E.g. rebalance after every 20 insertions ⟹ choose *k=20*

Note: To do this efficiently we would need to change tree data structure and basic operations:

```
typedef struct Node {
   int  data;
   int  nnodes;      // #nodes in my tree
   Tree left, right; // subtrees
} Node;
```

... Rebalancing Trees

How to rebalance a BST?   Move median item to root.

... Rebalancing Trees

Implementation of rebalance:

```
rebalance(t):
│   Input  tree t with n nodes
│   Output t rebalanced
│
│   if n≥3 then
│   │   t=partition(t,⌊n/2⌋)         // put node with median key at root
│   │   left(t)=rebalance(left(t))   // then rebalance each subtree
│   │   right(t)=rebalance(right(t))
│   end if
│   return t
```

New operation on trees:

- **partition(tree,i)**: re-arrange tree so that element with index i becomes root

For tree with N nodes, indices are $0 .. N\text{-}1$

Partition: moves i th node to root

Implementation of partition operation:

```
partition(tree,i):
│   Input  tree with n nodes, index i
│   Output tree with iᵗʰ item moved to the root
│
│   m=#nodes(left(tree))
│   if i < m then
│      left(tree)=partition(left(tree),i)
│      tree=rotateRight(tree)
│   else if i > m then
│      right(tree)=partition(right(tree),i-m-1)
│      tree=rotateLeft(tree)
```

```
  | end if
  | return tree
```

Note: size(tree) = n, size(left(tree)) = m, size(right(tree)) = n-m-1   (why -1?)

Exercise #8: Partition

Consider the tree `t`:

Show the result of `partition(t,3)`

... Rebalancing Trees

Analysis of rebalancing: visits every node $\Rightarrow O(N)$

Cost means not feasible to rebalance after each insertion.

When to rebalance? … Some possibilities:

- after every k insertions
- whenever "imbalance" exceeds threshold

Either way, we tolerate worse search performance for periods of time.

Does it solve the problem? … Not completely $\Rightarrow$ Solution: real balanced trees (next week)

Application of BSTs: Sets

Trees provide efficient search.

Sets require efficient search

- to find where to insert/delete
- to test for set membership

Logical to implement a set ADT via `BSTree`

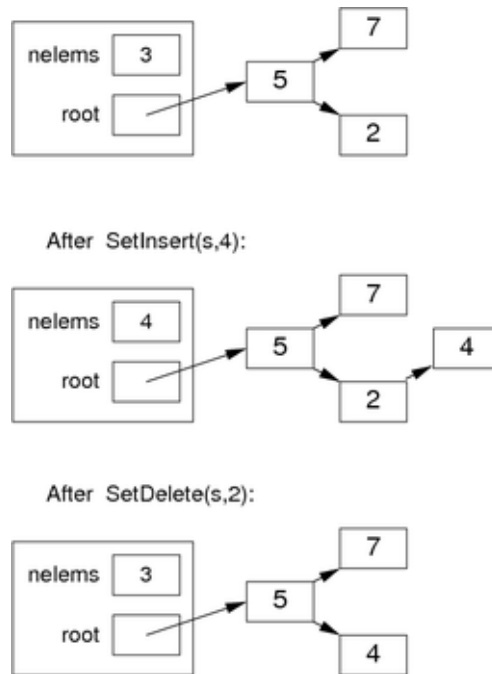## ... Application of BSTs: Sets

Assuming we have `Tree` implementation

- which precludes duplicate key values
- which implements insertion, search, deletion

then `Set` implementation is

- `SetInsert(Set,Item)` ≡ `TreeInsert(Tree,Item)`
- `SetDelete(Set,Item)` ≡ `TreeDelete(Tree,Item.Key)`
- `SetMember(Set,Item)` ≡ `TreeSearch(Tree,Item.Key)`

... Application of BSTs: Sets

... Application of BSTs: Sets

Concrete representation:

```c
#include <BSTree.h>

typedef struct SetRep {
    int    nelems;
    Tree   root;
} SetRep;

Set newSet() {
    Set S = malloc(sizeof(SetRep));
    assert(S != NULL);
    S->nelems = 0;
    S->root = newTree();
    return S;
}
```

Summary

- Binary search tree (BST) data structure
- BST insertion and deletion

- Other tree operations
  - tree rotation
  - tree partition
  - joining trees

- Suggested reading:
  - Sedgewick, Ch.12.5-12.6,12.8-12.9